

```
In [ ]: from keras.datasets import mnist
import matplotlib.pyplot as plt

data = mnist.load_data()


In [ ]: (X_train, y_train), (X_test, y_test) = data

In [ ]: X_train.shape

Out[ ]: (60000, 28, 28)

In [ ]: plt.figure(figsize=(10, 10))
for i in range(9):
    # Display the image
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(X_train[i], cmap='gray')
    # Print the label
    plt.title(y_train[i])
    plt.axis('off')

# Show the subplot
plt.show()
```



```
In [ ]: X_train = X_train.reshape((X_train.shape[0], 28*28)).astype('float32')
X_test = X_test.reshape((X_test.shape[0], 28*28)).astype('float32')

In [ ]: X_train = X_train / 255
X_test = X_test / 255

In [ ]: from keras import utils

print(y_test.shape)

y_train = utils.to_categorical(y_train)
y_test = utils.to_categorical(y_test)

num_classes = y_test.shape[1]
```

```
print(y_test.shape)
```

(10000,)
(10000, 10)

```
In [ ]: from keras.models import Sequential
        from keras.layers import Dense
```

```
In [ ]: model = Sequential()

        model.add(Dense(32, input_dim = 28 * 28, activation= 'relu'))

        model.add(Dense(64, activation = 'relu'))

        model.add(Dense(10, activation = 'softmax'))
```

```
In [ ]: model.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
```

```
In [ ]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| ===== | | |
| dense (Dense) | (None, 32) | 25120 |
| dense_1 (Dense) | (None, 64) | 2112 |
| ===== | | |
| Layer (type) | Output Shape | Param # |
| ===== | | |
| dense (Dense) | (None, 32) | 25120 |
| dense_1 (Dense) | (None, 64) | 2112 |
| dense_2 (Dense) | (None, 10) | 650 |
| ===== | | |
| Total params: 27882 (108.91 KB) | | |
| Trainable params: 27882 (108.91 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |
| ===== | | |

```
In [ ]: model.fit(X_train, y_train, epochs= 10, batch_size = 100)
```

Epoch 1/10
600/600 [=====] - 2s 2ms/step - loss: 0.4241 - accuracy: 0.8765
Epoch 2/10
600/600 [=====] - 1s 2ms/step - loss: 0.1880 - accuracy: 0.9458
Epoch 3/10
600/600 [=====] - 1s 2ms/step - loss: 0.1482 - accuracy: 0.9576
Epoch 4/10
600/600 [=====] - 1s 2ms/step - loss: 0.1238 - accuracy: 0.9635
Epoch 5/10
600/600 [=====] - 1s 1ms/step - loss: 0.1076 - accuracy: 0.9685
Epoch 6/10
600/600 [=====] - 1s 2ms/step - loss: 0.0952 - accuracy: 0.9716
Epoch 7/10
600/600 [=====] - 1s 2ms/step - loss: 0.0854 - accuracy: 0.9745
Epoch 8/10
600/600 [=====] - 1s 2ms/step - loss: 0.0757 - accuracy: 0.9772
Epoch 9/10
600/600 [=====] - 1s 2ms/step - loss: 0.0700 - accuracy: 0.9789
Epoch 10/10
600/600 [=====] - 1s 2ms/step - loss: 0.0639 - accuracy: 0.9804

Out[]: <keras.src.callbacks.History at 0x2a128de1f90>

```
In [ ]: scores = model.evaluate(X_test, y_test)
        print('Accuracy: ',scores[1] * 100)
```

313/313 [=====] - 0s 1ms/step - loss: 0.0993 - accuracy: 0.9694
Accuracy: 96.93999886512756